

# Are You Sure You Want to Use MMAP in Your Database Management System?

MMAP = ☺

Junpeng Zhu

Greenplum, VMware, Inc.

Apr, 28, 2022





- 1 Background
  - Recap
  - What's the Problems We Concern ?
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions
  - Authors Opinions
  - My Opinions
- 5 Acknowledgements and Questions

# Outline



- 1 Background
  - Recap
  - What's the Problems We Concern ?
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions
- 5 Acknowledgements and Questions



# Outline

- 1 Background
  - Recap
  - What's the Problems We Concern ?
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions
- 5 Acknowledgements and Questions

[Andrew Crotty](#)

Carnegie Mellon University

[Viktor Leis](#)

Friedrich-Alexander-Universität

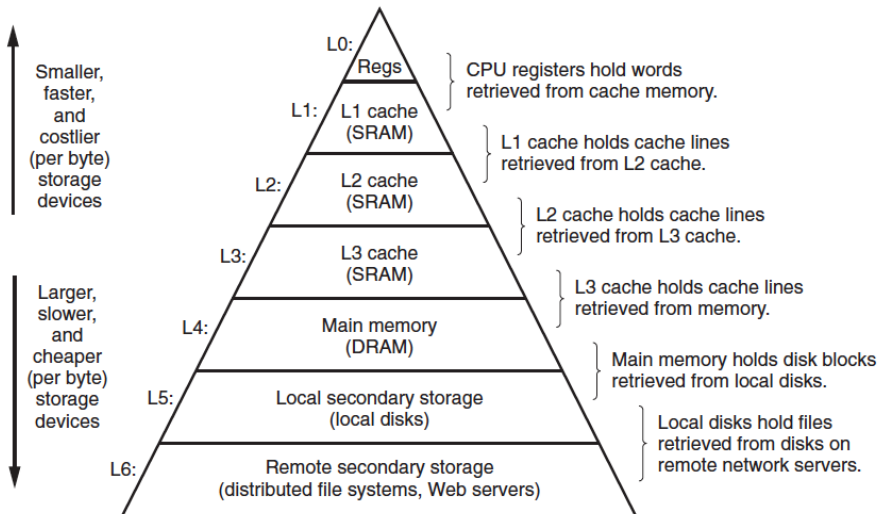
[Andy Pavlo](#)

Carnegie Mellon University

- Andrew Crotty
  - ✓ [Andrew Crotty's Bio](#)
  - ✓ Ph.D at Brown, in 2019. Post-doctoral at CMU.
- Viktor Leis
  - ✓ [Viktor Leis Bio](#)
  - ✓ Ph.D at TUM, full professor at Friedrich Schiller University Jena.
- Andy Pavlo
  - ✓ [Andy Pavlo Bio](#)
  - ✓ Ph.D at Brown, associate professor at CMU.
  - ✓ **Never use mmap in a DBMS at his tombstone.**




# Storage hierarchy, Cont.





# Architecture of RDBMS, Cont.

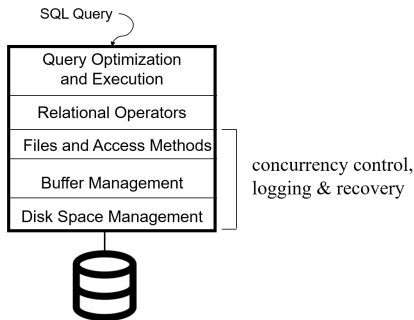
- Query optimization and execution
- Relational operators
- Files and access methods
- Buffer pool management
- Disk space management





# Architecture of RDBMS, Cont.

- Query optimization and execution
- Relational operators
- Files and access methods
- Buffer pool management
- Disk space management



- Crash recovery is awfully difficult !
  - \* The recovery system **depends on behavior of many other components** of DBMS, such as concurrency control, buffer management, disk management, and query processing.



# Buffer Pool Management, Cont.



GREENPLUM  
DATABASE®

- **Force policy** – make sure that every update is on the DB disk before **commit**.
  - ✓ Provides durability without REDO logging.
  - ✓ But, can cause poor performance **due to a large random write operations**.

# Buffer Pool Management, Cont.



GREENPLUM  
DATABASE®

- **Force policy** – make sure that every update is on the DB disk before **commit**.
  - ✓ Provides durability without REDO logging.
  - ✓ But, can cause poor performance **due to a large random write operations**.
- **No Steal policy** – don't allow buffer pool frames with uncommitted updates to overwrite committed data on DB disk.
  - ✓ Useful for ensuring atomicity without UNDO logging.
  - ✓ But can cause poor performance due to **(1) A larger buffer is required ; or (2) writing that data to temporary location on non-volatile storage (e.g., swap area)**.



# Buffer Pool Management, Cont.

- **Force policy** – make sure that every update is on the DB disk before **commit**.
  - ✓ Provides durability without REDO logging.
  - ✓ But, can cause poor performance **due to a large random write operations**.
- **No Steal policy** – don't allow buffer pool frames with uncommitted updates to overwrite committed data on DB disk.
  - ✓ Useful for ensuring atomicity without UNDO logging.
  - ✓ But can cause poor performance due to **(1) A larger buffer is required ; or (2) writing that data to temporary location on non-volatile storage (e.g., swap area)**.

**In practice, even to get Force/No-Steal to work requires some nasty details for handling unexpected failures. . .**

# Buffer Pool Replace Policy, Cont.



GREENPLUM  
DATABASE®

- **No Force**

- ▶ What if system crashes before a modified page written by a committed transaction makes it to DB disk?
  - ✓ Write as little as possible, in a convenient place, at commit time, to support REDOing modifications. → **WAL Logging**.

# Buffer Pool Replace Policy, Cont.



- **No Force**

- ▶ What if system crashes before a modified page written by a committed transaction makes it to DB disk?
  - ✓ Write as little as possible, in a convenient place, at commit time, to support REDOing modifications. → **WAL Logging**.

- **Steal**

- ▶ What if a transaction that performed updates aborts? → **WAL Logging**
- ▶ What if system crashes before transaction is finished? → **WAL Logging**
  - ✓ Must remember the old value of P (to support UNDOing the write to page P).



# Buffer Pool Replace Policy, Cont.

- No Force

- ▶ What if system crashes before a modified page written by a committed transaction makes it to DB disk?
  - ✓ Write as little as possible, in a convenient place, at commit time, to support REDOing modifications. → WAL Logging.

- Steal

- ▶ What if a transaction that performed updates aborts? → WAL Logging
- ▶ What if system crashes before transaction is finished? → WAL Logging
  - ✓ Must remember the old value of P (to support UNDOing the write to page P).

No force and steal buffer pool policy is commonly used in modern DBMSs.

# Buffer Pool Management, Cont.



GREENPLUM  
DATABASE®

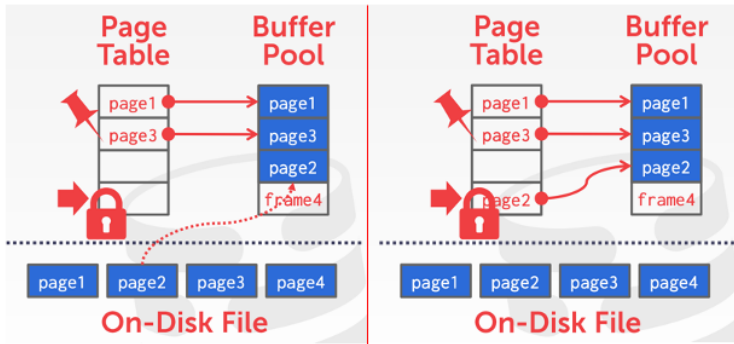
|          | No Steal       | Steal          |
|----------|----------------|----------------|
| No Force |                | <b>Fastest</b> |
| Force    | <b>Slowest</b> |                |

Performance Implications

|          | No Steal                   | Steal                   |
|----------|----------------------------|-------------------------|
| No Force | <b>No UNDO<br/>REDO</b>    | <b>UNDO<br/>REDO</b>    |
| Force    | <b>No UNDO<br/>No REDO</b> | <b>UNDO<br/>No REDO</b> |

Log/Recovery Implications

# Buffer Pool Management, Cont. <sup>1</sup>

GREENPLUM  
DATABASE®

## 1. CMU 15-445/645 Fall 2021 Buffer Pool Slide





# MMAP as Buffer Pool, Cont. <sup>2</sup>

- Memory-mapped (mmap) file I/O is an OS-provided feature.
  - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.

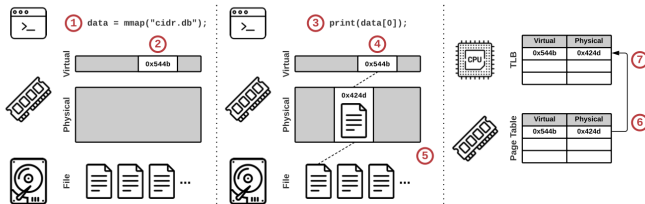
---

## 2. Ethanzjp MMAP Gendata Demo



# MMAP as Buffer Pool, Cont. <sup>2</sup>

- Memory-mapped (mmap) file I/O is an OS-provided feature.
  - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.
  - ✓ The program then accesses pages via **pointers** as if the file resided entirely in memory.



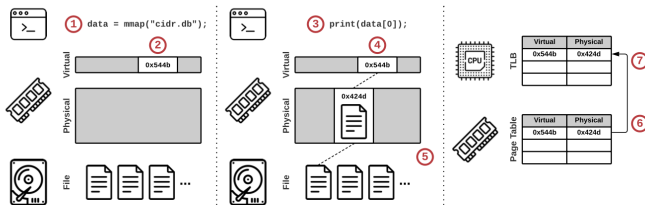
Step-by-step illustration of how a program accesses a file using mmap.

## 2. Ethanzjp MMAP Gendata Demo



# MMAP as Buffer Pool, Cont. <sup>2</sup>

- Memory-mapped (mmap) file I/O is an OS-provided feature.
  - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.
  - ✓ The program then accesses pages via **pointers** as if the file resided entirely in memory.
  - ✓ The OS **transparently loads pages** only when the program references them.



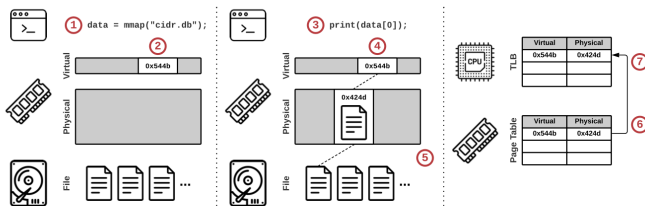
Step-by-step illustration of how a program accesses a file using mmap.

## 2. Ethanzjp MMAP Gendata Demo



# MMAP as Buffer Pool, Cont. <sup>2</sup>

- Memory-mapped (mmap) file I/O is an OS-provided feature.
  - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.
  - ✓ The program then accesses pages via **pointers** as if the file resided entirely in memory.
  - ✓ The OS **transparently loads pages** only when the program references them.
  - ✓ The OS **automatically evicts pages** if memory fills up.



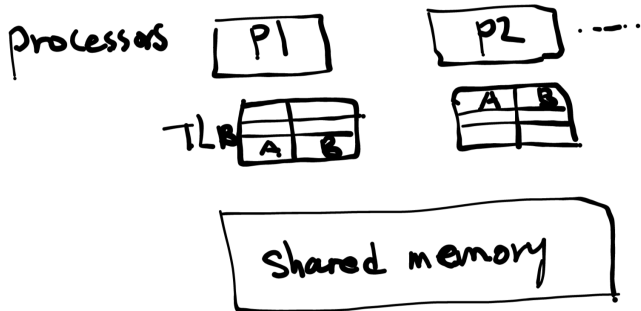
Step-by-step illustration of how a program accesses a file using mmap.

## 2. Ethanzjp MMAP Gendata Demo



# TLB shutdown

- Shared Memory Model<sup>3</sup>



## 3. Ethanzjp Shared Memory Demo

# Comparision of Buffer Pool and MMAP



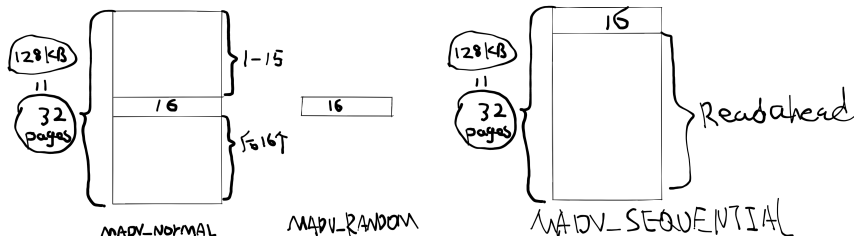
- Buffer Pool
  - ✓ The DBMS maintaining complete control over how and when it transfers pages.
- MMAP
  - ✓ The OS handles all necessary paging behind the scenes rather than the DBMS's buffer pool.
- Stonebraker 1981 opinion <sup>4</sup>

However, many DBMSs including INGRES [20] and System R [4] choose to put a DBMS managed buffer pool in user space to reduce overhead. Hence, each of these systems has gone to the trouble of constructing its own buffer pool manager to enhance performance.



# POSIX API, Cont.

- `mmap`<sup>5</sup>
- `madvise` hints to the OS about expected data access patterns<sup>6</sup>



- `mlock`<sup>7</sup> allows DBMS pin memory. But OS is permitted to flush dirty pages to the backing file at any time, even if the page is pinned.
- `msync`<sup>8</sup> explicitly flushes the specified memory range to secondary storage.

- 
5. `mmap` man7 page
  6. `madvise` man7 page
  7. `mlock` man7 page
  8. `msync` man7 page

# Modern MMAP-based DBMS



| DBMS        | MMAP Use  | Details  |
|-------------|-----------|----------|
| MonetDB     | 2002–     | [12, 21] |
| MongoDB     | 2009–2019 | [14, 3]  |
| LevelDB     | 2011–     | [5]      |
| LMDB        | 2011–     | [20]     |
| SQLite      | 2013–     | [7]      |
| SingleStore | 2013–2015 | [32]     |
| QuestDB     | 2014–     | [34]     |
| RavenDB     | 2014–     | [4]      |
| InfluxDB    | 2015–2020 | [8, 1]   |
| WiredTiger  | 2020–     | [17]     |

## 9. LevelDB Snapshot Demo



# Modern MMAP-based DBMS



| DBMS        | MMAP Use  | Details  |
|-------------|-----------|----------|
| MonetDB     | 2002–     | [12, 21] |
| MongoDB     | 2009–2019 | [14, 3]  |
| LevelDB     | 2011–     | [5]      |
| LMDB        | 2011–     | [20]     |
| SQLite      | 2013–     | [7]      |
| SingleStore | 2013–2015 | [32]     |
| QuestDB     | 2014–     | [34]     |
| RavenDB     | 2014–     | [4]      |
| InfluxDB    | 2015–2020 | [8, 1]   |
| WiredTiger  | 2020–     | [17]     |

- MonetDB use WiredTiger and MMAPv1 (optional) as storage engine.
- InfluxDB replaced mmap after observing I/O spikes for writes when a database grew larger than a few GB in size.

## 9. LevelDB Snapshot Demo

# Modern MMAP-based DBMS



| DBMS        | MMAP Use  | Details  |
|-------------|-----------|----------|
| MonetDB     | 2002–     | [12, 21] |
| MongoDB     | 2009–2019 | [14, 3]  |
| LevelDB     | 2011–     | [5]      |
| LMDB        | 2011–     | [20]     |
| SQLite      | 2013–     | [7]      |
| SingleStore | 2013–2015 | [32]     |
| QuestDB     | 2014–     | [34]     |
| RavenDB     | 2014–     | [4]      |
| InfluxDB    | 2015–2020 | [8, 1]   |
| WiredTiger  | 2020–     | [17]     |

- MonetDB use WiredTiger and MMAPv1 (optional) as storage engine.
- InfluxDB replaced mmap after observing I/O spikes for writes when a database grew larger than a few GB in size.
- SingleStore removed mmap-based file I/O after encountering poor performance on simple sequential scan queries.

## 9. LevelDB Snapshot Demo

# Modern MMAP-based DBMS



| DBMS        | MMAP Use  | Details  |
|-------------|-----------|----------|
| MonetDB     | 2002–     | [12, 21] |
| MongoDB     | 2009–2019 | [14, 3]  |
| LevelDB     | 2011–     | [5]      |
| LMDB        | 2011–     | [20]     |
| SQLite      | 2013–     | [7]      |
| SingleStore | 2013–2015 | [32]     |
| QuestDB     | 2014–     | [34]     |
| RavenDB     | 2014–     | [4]      |
| InfluxDB    | 2015–2020 | [8, 1]   |
| WiredTiger  | 2020–     | [17]     |

- MonetDB use WiredTiger and MMAPv1 (optional) as storage engine.
- InfluxDB replaced mmap after observing I/O spikes for writes when a database grew larger than a few GB in size.
- SingleStore removed mmap-based file I/O after encountering poor performance on simple sequential scan queries.
- RocksDB replace mmap as a fork of LevelDB<sup>9</sup>.

## 9. LevelDB Snapshot Demo



# What is the truth ?

The DBMS seems no longer needs to manage its own buffer pool, as it cedes this responsibility to the OS.



# Outline

- 1 Background
  - Recap
  - What's the Problems We Concern ?
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions
- 5 Acknowledgements and Questions

# When Building a Database, We Need<sup>10</sup>



GREENPLUM  
DATABASE®

- Reading the data from disk.
- Concurrency between different threads reading the same data.
- Caching and buffer management.
- Eviction of pages from memory.
- Playing nice with other processes in the machine.
- Tracking dirty pages and writing to disk.

---

10. CMU-15-445/645 Buffer Pool Management

# Outline



GREENPLUM  
DATABASE®

- 1 Background
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions
- 5 Acknowledgements and Questions

# Transactional Safety



- A database needs to know when the data is persisted to disk. When using mmap, we explicitly give up that knowledge.
  - \* Single writer and multiple readers are common in embedded systems, e.g. LMDB, LevelDB, Lucene.
  - \* The interleaving writes is a must for Postgres due that may have a transaction spanning multiple network calls.
- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
  - \* Due to transparent paging, **the OS can flush a dirty page (steal)** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
  - \* The DBMS cannot prevent these flushes and receives **no warning** when they occur.



# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW

# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.

# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.

# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.
    - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.

# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.
    - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
    - ✓ DBMS applies the committed changes to the primary copy using msync.

# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.
    - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
    - ✓ DBMS applies the committed changes to the primary copy using msync.
  - \* Common Used User space CoW

# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.
    - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
    - ✓ DBMS applies the committed changes to the primary copy using msync.
  - \* Common Used User space CoW
    - ✓ Copy the mmap-backed memory page to user buffer.

# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.
    - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
    - ✓ DBMS applies the committed changes to the primary copy using msync.
  - \* Common Used User space CoW
    - ✓ Copy the mmap-backed memory page to user buffer.
    - ✓ Update and recording the WAL logging.



# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.
    - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
    - ✓ DBMS applies the committed changes to the primary copy using msync.
  - \* Common Used User space CoW
    - ✓ Copy the mmap-backed memory page to user buffer.
    - ✓ Update and recording the WAL logging.
    - ✓ Copy the modified pages back to the mmap-backed memory. (When do the DBMSs copy user-space content to mmap-backed memory ? **Concurrent transactions are all end !**)


# Handling Updates When it May Use MMAP

GREENPLUM  
DATABASE®

- Three categories for handling updates :
  - \* OS CoW
    - ✓ MAP\_PRIVATE to enable OS CoW.
    - ✓ The DBMS modifies the affected pages in the private workspace.
    - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
    - ✓ DBMS applies the committed changes to the primary copy using msync.
  - \* Common Used User space CoW
    - ✓ Copy the mmap-backed memory page to user buffer.
    - ✓ Update and recording the WAL logging.
    - ✓ Copy the modified pages back to the mmap-backed memory. (When do the DBMSs copy user-space content to mmap-backed memory? **Concurrent transactions are all end !**)
  - \* Shadow Paging



# System R's Shadow paging



- Master : Contains only changes from committed txns.
- Shadow : Temporary db with changes made from uncommitted txns.

# IO Stalls



- Accessing any page could result in an unexpected I/O stall because the DBMS cannot know whether the page is in memory.
  - ✓ Pinning memory.
  - ✓ mlock the memory.
  - ✓ madvise, but os is free ignore the advise.

# Error Handling



- Page-level checksums
- Gracefully handling I/O errors
  - \* Error return value for system calls.
  - \* SIGBUS.



# Performance Issues

- Page table contention
  - \* PTE and PMD split spin\_lock.<sup>11</sup>
- Single-threaded page eviction and for larger-than- memory DBMS workloads on high-bandwidth secondary storage devices.
- TLB shootdowns
  - \* Fast I/O
  - \* Working set that significantly exceeds memory
  - \* No other work that needs to be done for processing a request

---

## 11. Page Table Split Lock

# Outline




GREENPLUM  
DATABASE®

- 1 Background
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis**
- 4 Conclusions
- 5 Acknowledgements and Questions



# Random Reads on Bandwidth<sup>12</sup>



(a) Bandwidth

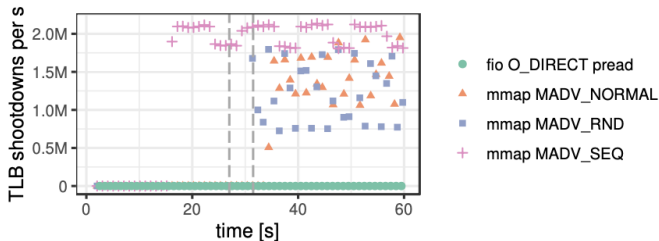
- Random access pattern over a 2 TB SSD range to simulate a larger-than-memory OLTP workload.
- The page cache had only 100 GB of memory, 95% of all accesses resulted in page faults
- fio baseline exhibited stable performance and achieved close to 900K reads per second

## 12. mmapbenchmark





# Random Reads on TLB Shutdown




**(b)** TLB Shutdowns


- we measured using `/proc/interrupt`



# Sequential Scan on Bandwidth



**Figure 3:** Sequential Scan – 1 SSD (mmap: 20 threads; fio: libaio, 1 thread, iodepth 256)



**Figure 4:** Sequential Scan – 10 SSDs (mmap: 20 threads; fio: libaio, 4 threads, iodepth 256)

# Outline



GREENPLUM  
DATABASE®

- 1 Background
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions**
  - Authors Opinions
  - My Opinions
- 5 Acknowledgements and Questions

# Outline



GREENPLUM  
DATABASE®

- 1 Background
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions**
  - Authors Opinions
  - My Opinions
- 5 Acknowledgements and Questions

# MMAP is Nerver Used!!!



GREENPLUM  
DATABASE®

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.

# MMAP is Nerver Used!!!



GREENPLUM  
DATABASE®

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.
  - \* You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.

# MMAP is Nerver Used!!!



GREENPLUM  
DATABASE®

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.
  - \* You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
  - \* You care about error handling and need to return correct results.

# MMAP is Never Used!!!



GREENPLUM  
DATABASE®

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.
  - \* You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
  - \* You care about error handling and need to return correct results.
  - \* You require high throughput on fast persistent storage devices.



# MMAP is Nerver Used !!!



GREENPLUM  
DATABASE®

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.
  - \* You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
  - \* You care about error handling and need to return correct results.
  - \* You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS :

# MMAP is Never Used!!!



GREENPLUM  
DATABASE®

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.
  - \* You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
  - \* You care about error handling and need to return correct results.
  - \* You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS :
  - ✓ Your working set (or the entire database) fits in memory and the workload is read-only.



# MMAP is Never Used!!!

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.
  - \* You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
  - \* You care about error handling and need to return correct results.
  - \* You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS :
  - ✓ Your working set (or the entire database) fits in memory and the workload is read-only.
  - ✓ You need to rush a product to the market and do not care about data consistency or long-term engineering headaches.



# MMAP is Nerver Used!!!

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS :
  - \* You need to perform updates in a transactionally safe fashion.
  - \* You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
  - \* You care about error handling and need to return correct results.
  - \* You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS :
  - ✓ Your working set (or the entire database) fits in memory and the workload is read-only.
  - ✓ You need to rush a product to the market and do not care about data consistency or long-term engineering headaches.
  - ✓ Otherwise, never.

# Outline



GREENPLUM  
DATABASE®

- 1 Background
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions**
  - Authors Opinions
  - My Opinions
- 5 Acknowledgements and Questions

# Unfair Comparision



- The problem exists, but the reasons and experiments are not sufficient.
- There are also many problems with buffer pool management.

# Outline



GREENPLUM  
DATABASE®

- 1 Background
- 2 Problem with MMAP : The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions
- 5 Acknowledgements and Questions**

Thank you !  
Welcome for any questions !



Junpeng Zhu  
Greenplum, VMware, Inc.