

Are You Sure You Want to Use MMAP in Your Database Management System?

MMAP = 😊

Junpeng Zhu

Greenplum, VMware, Inc.
zjunpeng@vmware.com

2022 of Greenplum TechTalk
May 3, 2022



**GREENPLUM
DATABASE®**

- 1 Background
- 2 Problem with MMAP: The Four Deadly Sins
- 3 Experimental Analysis
- 4 Conclusions
- 5 References
- 6 Acknowledgements and Questions



[Andrew Crotty](#)

Carnegie Mellon University



[Viktor Leis](#)

Friedrich-Alexander-Universität




[Andy Pavlo](#)

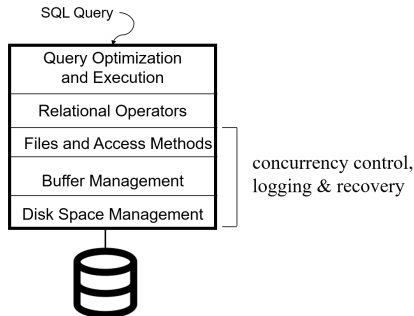
Carnegie Mellon University

- Andrew Crotty
 - ✓ [Andrew Crotty's Bio](#)
 - ✓ Ph.D at Brown, in 2019. Post-doctoral at CMU.
- Viktor Leis
 - ✓ [Viktor Leis Bio](#)
 - ✓ Ph.D at TUM, full professor at Friedrich Schiller University Jena.
- Andy Pavlo
 - ✓ [Andy Pavlo Bio](#)
 - ✓ Ph.D at Brown, associate professor at CMU.
 - ✓ **Never use mmap in a DBMS at his tombstone.**


Storage hierarchy, Cont.



- Query optimization and execution
- Relational operators
- Files and access methods
- Buffer pool management
- Disk space management



- Query optimization and execution
- Relational operators
- Files and access methods
- Buffer pool management
- Disk space management



- Crash recovery is awfully difficult!
 - * The recovery system **depends on behavior of many other components** of DBMS, such as concurrency control, buffer management, disk management, and query processing.

- **Force policy** – make sure that every update is on the DB disk before **commit**.
 - ✓ Provides durability without REDO logging.
 - ✓ But, can cause poor performance **due to a large random write operations**.

- **Force policy** – make sure that every update is on the DB disk before **commit**.
 - ✓ Provides durability without REDO logging.
 - ✓ But, can cause poor performance **due to a large random write operations**.
- **No Steal policy** – don't allow buffer pool frames with uncommitted updates to overwrite committed data on DB disk.
 - ✓ Useful for ensuring atomicity without UNDO logging.
 - ✓ But can cause poor performance due to (1) **A larger buffer is required**; or (2) **writing that data to temporary location on non-volatile storage (e.g., swap area)**.

- **Force policy** – make sure that every update is on the DB disk before **commit**.
 - ✓ Provides durability without REDO logging.
 - ✓ But, can cause poor performance **due to a large random write operations**.
- **No Steal policy** – don't allow buffer pool frames with uncommitted updates to overwrite committed data on DB disk.
 - ✓ Useful for ensuring atomicity without UNDO logging.
 - ✓ But can cause poor performance due to (1) **A larger buffer is required**; or (2) **writing that data to temporary location on non-volatile storage (e.g., swap area)**.

In practice, even to get Force/No-Steal to work requires some nasty details for handling unexpected failures...

- **No Force**
 - What if system crashes before a modified page written by a committed transaction makes it to DB disk?
 - ✓ Write as little as possible, in a convenient place, at commit time, to support REDOing modifications. → **WAL Logging**.

- **No Force**

- What if system crashes before a modified page written by a committed transaction makes it to DB disk?
 - ✓ Write as little as possible, in a convenient place, at commit time, to support REDOing modifications. → **WAL Logging**.

- **Steal**

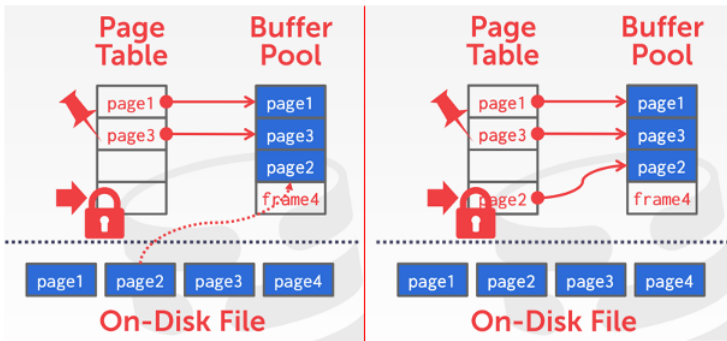
- What if a transaction that performed updates aborts? → **WAL Logging**
- What if system crashes before transaction is finished? → **WAL Logging**
 - ✓ Must remember the old value of P (to support UNDOing the write to page P).

	No Steal	Steal
No Force		Fastest
Force	Slowest	

Performance Implications

	No Steal	Steal
No Force	No UNDO REDO	UNDO REDO
Force	No UNDO No REDO	UNDO No REDO

Log/Recovery Implications

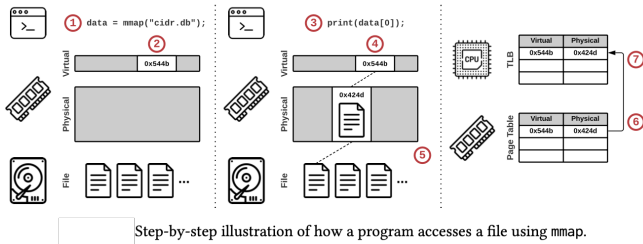


¹CMU 15-445/645 Fall 2021 Buffer Pool Slide

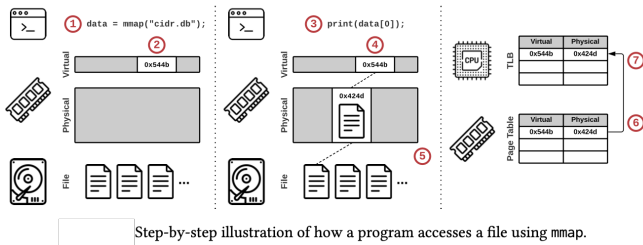
- Memory-mapped (mmap) file I/O is an OS-provided feature.
 - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.

²Ethanzjp MMAP Gendata Demo

- Memory-mapped (mmap) file I/O is an OS-provided feature.
 - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.
 - ✓ The program then accesses pages via **pointers** as if the file resided entirely in memory.




- Memory-mapped (mmap) file I/O is an OS-provided feature.
 - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.
 - ✓ The program then accesses pages via **pointers** as if the file resided entirely in memory.
 - ✓ The OS **transparently loads pages** only when the program references them.




Step-by-step illustration of how a program accesses a file using mmap.

- Memory-mapped (mmap) file I/O is an OS-provided feature.
 - ✓ It maps **the contents of a file** on secondary storage into a **program's virtual address space**.
 - ✓ The program then accesses pages via **pointers** as if the file resided entirely in memory.
 - ✓ The OS **transparently loads pages** only when the program references them.
 - ✓ The OS **automatically evicts pages** if memory fills up.



- Shared Memory Model³




³Ethanzjp Shared Memory Demo

- Buffer Pool
 - ✓ The DBMS maintaining complete control over how and when it transfers pages.
- MMAP
 - ✓ The OS handles all necessary paging behind the scenes rather than the DBMS's buffer pool.
- Stonebraker 1981 opinion ⁴

However, many DBMSs including INGRES [20] and System R [4] choose to put a DBMS managed buffer pool in user space to reduce overhead. Hence, each of these systems has gone to the trouble of constructing its own buffer pool manager to enhance performance.

⁴1981 Stonebraker's Paper

- mmap⁵
- madvise hints to the OS about expected data access patterns⁶



- mlock⁷ allows DBMS pin memory. But OS is permitted to flush dirty pages to the backing file at any time, even if the page is pinned.
- msync⁸ explicitly flushes the specified memory range to secondary storage.

⁵mmap man7 page

⁶madvise man7 page

⁷mlock man7 page

⁸msync man7 page

DBMS	MMAP Use	Details
MonetDB	2002–	[12, 21]
MongoDB	2009–2019	[14, 3]
LevelDB	2011–	[5]
LMDB	2011–	[20]
SQLite	2013–	[7]
SingleStore	2013–2015	[32]
QuestDB	2014–	[34]
RavenDB	2014–	[4]
InfluxDB	2015–2020	[8, 1]
WiredTiger	2020–	[17]

DBMS	MMAP Use	Details
MonetDB	2002–	[12, 21]
MongoDB	2009–2019	[14, 3]
LevelDB	2011–	[5]
LMDB	2011–	[20]
SQLite	2013–	[7]
SingleStore	2013–2015	[32]
QuestDB	2014–	[34]
RavenDB	2014–	[4]
InfluxDB	2015–2020	[8, 1]
WiredTiger	2020–	[17]

- MonetDB use WiredTiger and MMAPv1 (optional) as storage engine.
- InfluxDB replaced mmap after observing I/O spikes for writes when a database grew larger than a few GB in size.

⁹LevelDB Snapshot Demo

DBMS	MMAP Use	Details
MonetDB	2002–	[12, 21]
MongoDB	2009–2019	[14, 3]
LevelDB	2011–	[5]
LMDB	2011–	[20]
SQLite	2013–	[7]
SingleStore	2013–2015	[32]
QuestDB	2014–	[34]
RavenDB	2014–	[4]
InfluxDB	2015–2020	[8, 1]
WiredTiger	2020–	[17]

- MonetDB use WiredTiger and MMAPv1 (optional) as storage engine.
- InfluxDB replaced mmap after observing I/O spikes for writes when a database grew larger than a few GB in size.
- SingleStore removed mmap-based file I/O after encountering poor performance on simple sequential scan queries.

⁹LevelDB Snapshot Demo

DBMS	MMAP Use	Details
MonetDB	2002–	[12, 21]
MongoDB	2009–2019	[14, 3]
LevelDB	2011–	[5]
LMDB	2011–	[20]
SQLite	2013–	[7]
SingleStore	2013–2015	[32]
QuestDB	2014–	[34]
RavenDB	2014–	[4]
InfluxDB	2015–2020	[8, 1]
WiredTiger	2020–	[17]

- MonetDB use WiredTiger and MMAPv1 (optional) as storage engine.
- InfluxDB replaced mmap after observing I/O spikes for writes when a database grew larger than a few GB in size.
- SingleStore removed mmap-based file I/O after encountering poor performance on simple sequential scan queries.
- RocksDB replace mmap as a fork of LevelDB ⁹.

⁹LevelDB Snapshot Demo

The DBMS seems no longer needs to manage its own buffer pool, as it cedes this responsibility to the OS.

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.
 - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.
 - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
 - ✓ DBMS applies the committed changes to the primary copy using msync.

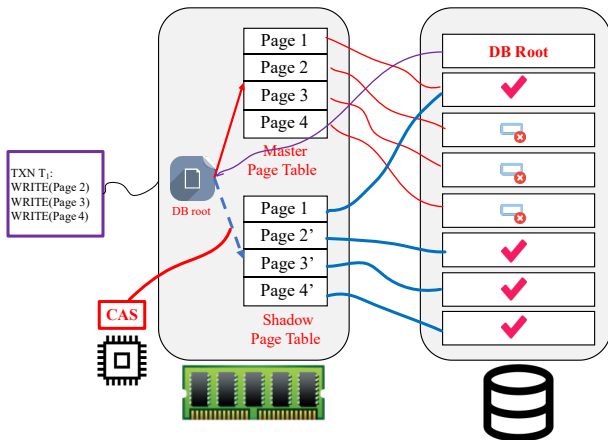
- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.
 - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
 - ✓ DBMS applies the committed changes to the primary copy using msync.
 - * User space CoW

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.
 - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
 - ✓ DBMS applies the committed changes to the primary copy using msync.
 - * User space CoW
 - ✓ Copy the mmap-backed memory page to user buffer.

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.
 - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
 - ✓ DBMS applies the committed changes to the primary copy using msync.
 - * User space CoW
 - ✓ Copy the mmap-backed memory page to user buffer.
 - ✓ Update and recording the WAL logging.

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.
 - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
 - ✓ DBMS applies the committed changes to the primary copy using msync.
 - * User space CoW
 - ✓ Copy the mmap-backed memory page to user buffer.
 - ✓ Update and recording the WAL logging.
 - ✓ Copy the modified pages back to the mmap-backed memory.

- The challenges inherent with guaranteeing transactional safety of modified pages in mmap-based DBMSs are well-known.
 - * Due to transparent paging, **the OS can flush a dirty page** to secondary storage **at any time**, irrespective of whether the writing transaction has committed.
 - * The DBMS cannot prevent these flushes and receives **no warning** when they occur.
- Three categories for handling updates:
 - * OS CoW
 - ✓ MAP_PRIVATE to enable OS CoW.
 - ✓ The DBMS modifies the affected pages in the private workspace.
 - ✓ To provide durability, the DBMS must use a write-ahead log (WAL) to record changes.
 - ✓ DBMS applies the committed changes to the primary copy using msync.
 - * User space CoW
 - ✓ Copy the mmap-backed memory page to user buffer.
 - ✓ Update and recording the WAL logging.
 - ✓ Copy the modified pages back to the mmap-backed memory.
 - * Shadow Paging




- Master: Contains only changes from committed txns.
- Shadow: Temporary db with changes made from uncommitted txns.

- Accessing any page could result in an unexpected I/O stall because the DBMS cannot know whether the page is in memory.
 - ✓ Pinning memory.
 - ✓ mlock the memory.
 - ✓ madvise, but os is free ignore the advise.

- page-level checksums
- gracefully handling I/O errors


- page table contention
- single-threaded page eviction and for larger-than- memory DBMS workloads on high-bandwidth secondary storage devices.
- TLB shutdowns.



(a) Bandwidth

- Random access pattern over a 2 TB SSD range to simulate a larger-than-memory OLTP workload.
- The page cache had only 100 GB of memory, 95% of all accesses resulted in page faults
- fio baseline exhibited stable performance and achieved close to 900K reads per second

¹⁰mmapbenchmark



(b) TLB Shootdowns

- we measured using /proc/interrupt




Figure 3: Sequential Scan – 1 SSD (mmap: 20 threads; fio: libaio, 1 thread, iodepth 256)




Figure 4: Sequential Scan – 10 SSDs (mmap: 20 threads; fio: libaio, 4 threads, iodepth 256)

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.
 - * You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.
 - * You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
 - * You care about error handling and need to return correct results.

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.
 - * You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
 - * You care about error handling and need to return correct results.
 - * You require high throughput on fast persistent storage devices.

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.
 - * You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
 - * You care about error handling and need to return correct results.
 - * You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS:

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.
 - * You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
 - * You care about error handling and need to return correct results.
 - * You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS:
 - ✓ Your working set (or the entire database) fits in memory and the workload is read-only.

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.
 - * You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
 - * You care about error handling and need to return correct results.
 - * You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS:
 - ✓ Your working set (or the entire database) fits in memory and the workload is read-only.
 - ✓ You need to rush a product to the market and do not care about data consistency or long-term engineering headaches.

- mmap is not a suitable replacement for a traditional buffer pool.
- When you should not use mmap in your DBMS:
 - * You need to perform updates in a transactionally safe fashion.
 - * You want to handle page faults without blocking on slow I/O or need explicit control over what data is in memory.
 - * You care about error handling and need to return correct results.
 - * You require high throughput on fast persistent storage devices.
- When you should maybe use mmap in your DBMS:
 - ✓ Your working set (or the entire database) fits in memory and the workload is read-only.
 - ✓ You need to rush a product to the market and do not care about data consistency or long-term engineering headaches.
 - ✓ Otherwise, never.



RavenDB Response



Community Comments

Thank you!
Welcome for any questions!



Junpeng Zhu
Greenplum, VMware, Inc.